

Machine Learning 1

Shivani Khosla (PID: A59010433)


10/22/2021

Clustering methods

Kmeans clustering in R is done with the 'kmeans()' function. Here we make up some data to test and learn with.

rev reverses order of vector cbind concatenates the two vectors in tmp

```
tmp <- c(rnorm(30, 3), rnorm(30, -3))  
data <- cbind(x = tmp, y = rev(tmp))  
plot(data)
```



Run 'kmeans()' set k to 2, nstart to 20. Have to tell it how many clusters you want.

```
km <- kmeans(data, centers = 2, nstart = 20)
km
```

[illegible]

Q. How many points are in each cluster?

km\$size

```
## [1] 30 30
```

Q. What ‘component’ of your result object details cluster assignment/membership?

```
km$cluster
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q. What ‘component’ of your result object details cluster center?

km\$centers

```
##           x           y
## 1  3.173994 -2.663504
## 2 -2.663504  3.173994
```

Q. Plot `x` colored by the `kmeans` cluster assignment and add cluster centers as blue points

```
plot(data, col = km$cluster)
points(km$centers, col = 'blue', pch = 15, cex = 2)
```



Hierarchical Clustering

Use the 'hclust()' function on the same data as before


```
hc <- hclust(dist(data))
hc
```

```
##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

hclust has a plot method

```
plot(hc)
```

Cluster Dendrogram



```
dist(data)
hclust (*, "complete")
```

To find membership vector, need to ‘cut’ the tree by using ‘cutree’ function and tell it the height to cut at


```
cutree(hc, h = 7)
```

[illegible]

Use 'cutree()' and state the number of k clusters we want

```
grps <- cutree(hc, k = 2)
```

```
plot(data, col = grps)
```



Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

How many rows and cols?

```
dim(x)
```

```
## [1] 17  5
```

```
x[, -1]
```

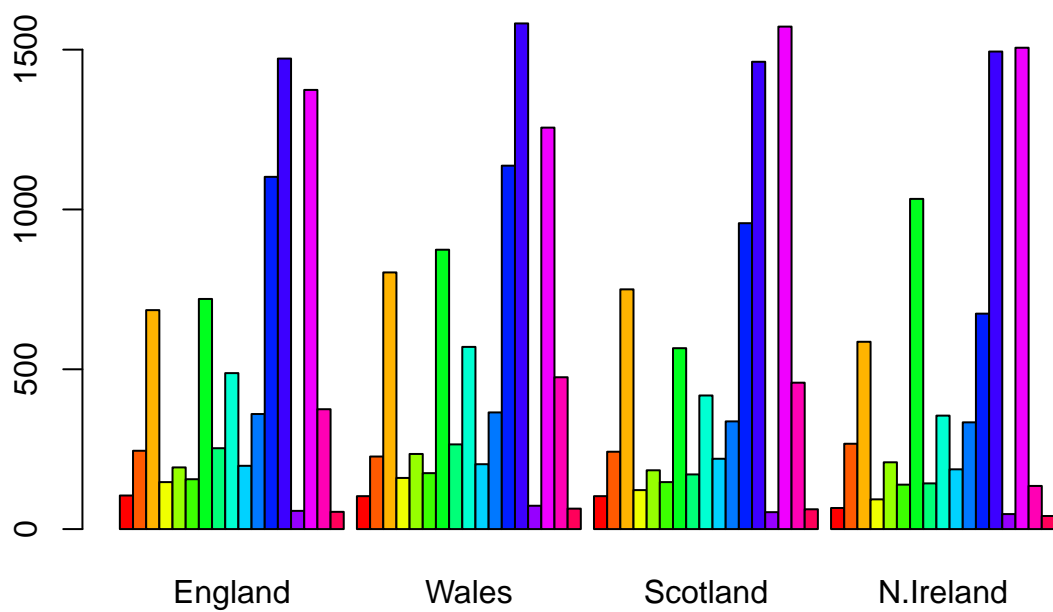
```
##      England Wales Scotland N.Ireland
## 1      105    103      103         66
## 2      245    227      242        267
## 3      685    803      750        586
## 4      147    160      122         93
## 5      193    235      184        209
## 6      156    175      147        139
## 7      720    874      566       1033
## 8      253    265      171        143
```

```
## 9      488  570      418      355
## 10     198  203      220      187
## 11     360  365      337      334
## 12    1102 1137      957      674
## 13    1472 1582     1462     1494
## 14      57   73       53       47
## 15    1374 1256     1572     1506
## 16     375  475      458      135
## 17      54   64       62       41
```

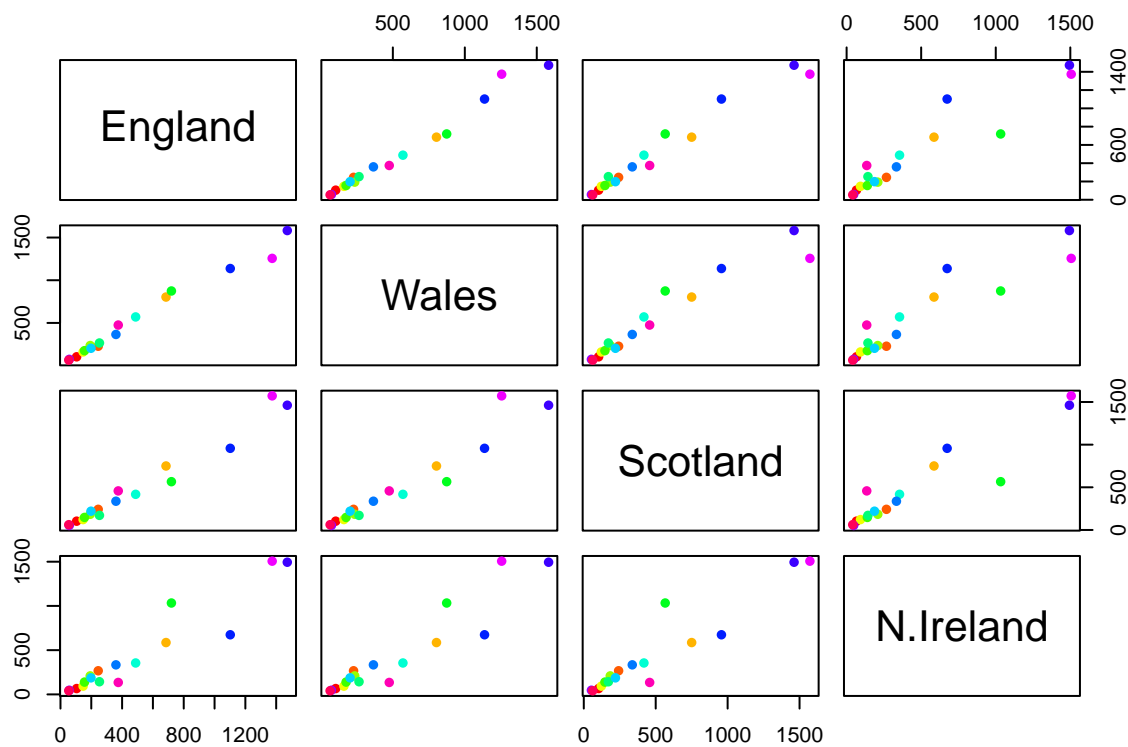
```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

```
##              England Wales Scotland N.Ireland
## Cheese              105   103      103       66
## Carcass_meat        245   227      242      267
## Other_meat          685   803      750      586
## Fish                147   160      122       93
## Fats_and_oils       193   235      184      209
## Sugars              156   175      147      139
## Fresh_potatoes      720   874      566     1033
## Fresh_Veg           253   265      171      143
## Other_Veg           488   570      418      355
## Processed_potatoes  198   203      220      187
## Processed_Veg       360   365      337      334
## Fresh_fruit         1102  1137      957      674
## Cereals             1472  1582     1462     1494
## Beverages           57    73       53       47
## Soft_drinks         1374  1256     1572     1506
## Alcoholic_drinks    375   475      458      135
## Confectionery        54    64       62       41
```

```
barplot(as.matrix(x), col = rainbow(17), beside = TRUE)
```



```
mycols <- rainbow(nrow(x))  
pairs(x, col = mycols, pch = 16)
```



PCA to the rescue!


Use base R function for PCA, which is called 'prcomp()'. This function wants the transpose of our data.

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

##	PC1	PC2	PC3	PC4
## Standard deviation	324.1502	212.7478	73.87622	4.189e-14
## Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
## Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
plot(pca)
```

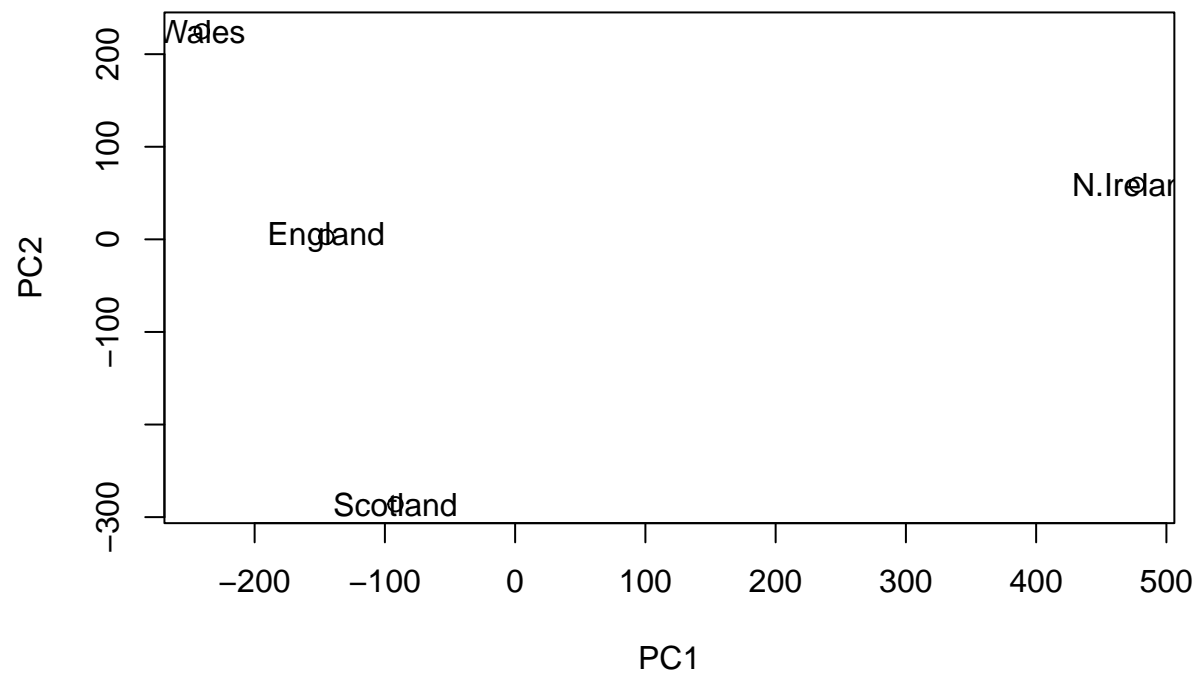
We want score plot (aka PCA plot). Basically of PC1 vs PC2

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation" "center"   "scale"    "x"
##
## $class
## [1] "prcomp"
```

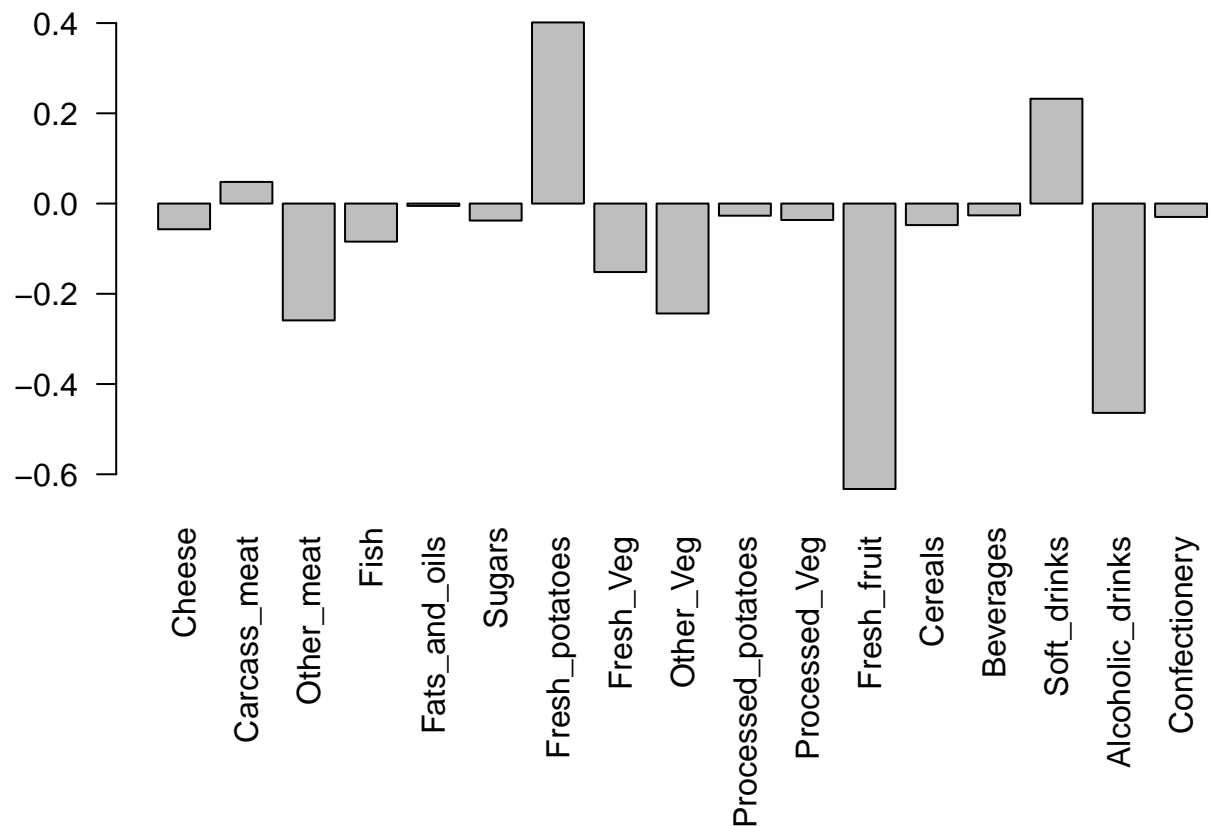
We are after the `pca$x` component for this plot

```
plot(pca$x[,1:2])
text(pca$x[,1:2], labels = colnames(x))
```



We can also examine the PCA “loadings”, which tell us how much the original variables contribute to each PC

```
par(mar=c(10, 3, 0.35, 0))  
barplot(pca$rotation[,1], las = 2)
```



One more PCA for today

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458  408  429 420  90  88  86  90  93
## gene2 219 200  204  210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4  783 792  829  856 760 849 856 835 885 894
## gene5  181 249  204  244 225 277 305 272 270 279
## gene6  460 502  491  491 493 612 594 577 618 638
```

```
nrow(rna.data)
```

```
## [1] 100
```

```
ncol(rna.data)
```

```
## [1] 10
```

```
colnames(rna.data)
```

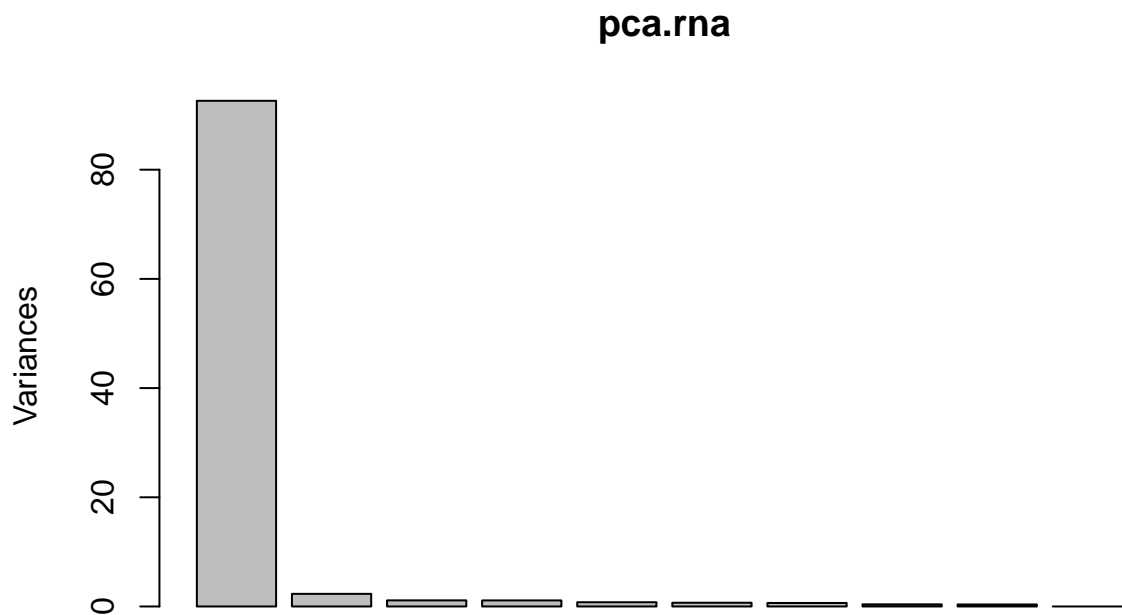
```
## [1] "wt1" "wt2" "wt3" "wt4" "wt5" "ko1" "ko2" "ko3" "ko4" "ko5"
```

```
pca.rna <- prcomp(t(rna.data), scale = TRUE)  
summary(pca.rna)
```

```
## Importance of components:
```

```
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7  
## Standard deviation  9.6237 1.5198 1.05787 1.05203 0.88062 0.82545 0.80111  
## Proportion of Variance 0.9262 0.0231 0.01119 0.01107 0.00775 0.00681 0.00642  
## Cumulative Proportion 0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251  
##              PC8      PC9      PC10  
## Standard deviation  0.62065 0.60342 3.348e-15  
## Proportion of Variance 0.00385 0.00364 0.000e+00  
## Cumulative Proportion 0.99636 1.00000 1.000e+00
```

```
plot(pca.rna)
```



```
plot(pca.rna$x[,1:2])  
text(pca.rna$x[,1:2], labels = colnames(rna.data))
```

